

Ch.4 Greedy Algorithms

Ex. (failure case)

coins 1, 10, 25, 100

use min # of coins to reach target
largest coin first (greedy strategy)

$T = 34$: $25 + 1 + \dots + 1$ (10 coins)

$10 + \dots + 10 + 1 + \dots + 1$ (7 coins)

4.1 Scheduling

Class 1, 2, ..., n, each class has starting time $s(n)$ and finishing time $f(n)$

Goal: schedule max # of classes that are compatible
(no overlap)

Idea: schedule classes 1 by 1 according to some ordering

	s	f	
a	0	4	a, b, c, d (sort by starting time)
b	1	3	solution: a → not always optimal
c	2	4	
d	3	5	d, a, c, b (sort by finishing time)

Greedy alg. O(nlogn)

sort by finishing time $T_{v_1}, T_{v_2}, \dots, T_{v_n}$, $f(T_{v_1}) \leq \dots \leq f(T_{v_n})$

$A = \emptyset$

for $i = 1, 2, \dots, n$:

if T_{v_i} is compatible w/ A :

$A \leftarrow A + [T_{v_i}]$

Thm. Greedy algorithm gives the optimal solution.

Pf. Let $\{i_1, i_2, \dots, i_k\}$ be the solution by greedy algorithm.

Claim: For any compatible classes $\{j_1, \dots, j_k\}$,

$$f(i_r) \leq f(j_r) \quad \forall r = 1, \dots, k$$

BC $r = 1$: $f(i_1) \leq f(j_1)$

IH Assume $f(i_r) \leq f(j_r) \quad \forall 1 \leq r \leq n$

IS WTS: $f(i_{n+1}) \leq f(j_{n+1})$

$$f(i_n) \leq f(j_n) \leq s(j_{n+1})$$

\Rightarrow class j_{n+1} is compatible with $\{i_1, i_2, \dots, i_m\}$

\Rightarrow GD will pick $i_{n+1} \Rightarrow f(i_{n+1}) \leq f(j_{n+1})$

If \exists a compatible set j_1, \dots, j_k s.t.

$f(i_k) \leq f(j_k) \leq s(j_{k+1})$, GD will be able to pick j_{k+1} .

Interval partition

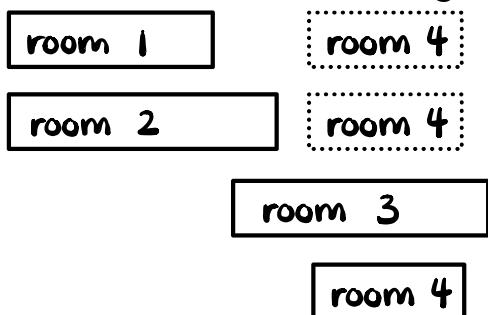
class 1, ..., n

Goal: min # of rooms to schedule all the classes

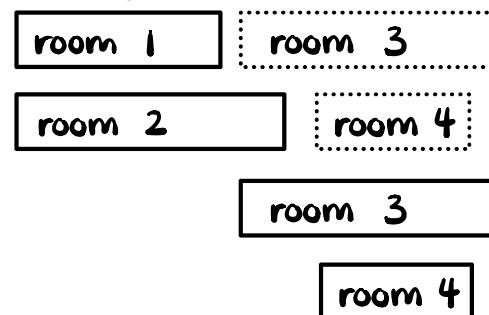
Idea: consider classes 1 by 1

create a room when needed

Sorting by finishing time:



Sorting by starting time:



Greedy alg.

sort classes by starting time: $S(\pi_1) \leq \dots \leq S(\pi_n)$

$d = 0$ # of rooms

for $j = 1, 2, \dots, n$:

if π_j is compatible with some room $k \leq d$

$$A_k = A_k + \{\pi_j\}$$

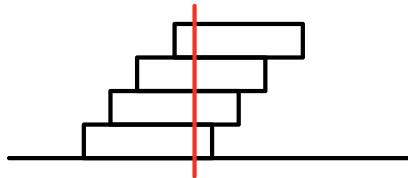
else:

create a new room $A_{d+1} = \{\pi_j\}$

$$d = d + 1$$

Pf. of correctness

"depth": max # of classes at any time



depth: lower bound of # rooms

Thm. Greedy outputs $d = \text{depth}$.

Pf. Assume room d is created by GD alg. when scheduling π_j .

$\Rightarrow \pi_j$ is not compatible with room $1, \dots, d-1$

Since we sort by starting time,

$$S(\pi_j) \geq S(\pi_i) \quad \forall i = 1, \dots, d-1.$$

π_j overlap with $x_i \Rightarrow f(x_i) > S(\pi_i)$

$\Rightarrow \forall i, S(\pi_i) \leq S(\pi_j) < f(x_i)$

$\Rightarrow S(\pi_j)$ has depth d .

4.4 Shortest Path

- Single source shortest path

- $G = (V, E)$ DG

- $l_e (e \in E)$: cost (length of each edge)

Goal: Find shortest path from a source node s to all other nodes.

- X : explored nodes (shortest path known)

- $d[u]$: min length $s \sim u$ path via X

Dijkstra's algorithm:

$$X = \{s\} \quad d[u] = \begin{cases} l(s, u) & \text{if } (s, u) \in E \\ \infty & \text{otherwise} \end{cases}$$

for $i = 1, \dots, n$:

$u \leftarrow \min d[u]$ among unexplored nodes

for each v s.t. $(u, v) \in E$

if $d[u] + l(u, v) < d[v]$

$d[v] = d[u] + l(u, v)$

$\text{pre}[v] = u$

shortest path $s \sim v$
" "
 $s \sim v = \text{pre}[v] - v$

print shortest path $s \sim v$ (backward):

$u = v$

while $u \neq s$:

print u

$u = \text{pre}[u]$

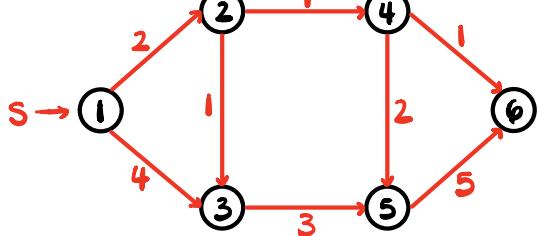
① unexplored node v with min $d[u]$

$\Rightarrow d[u]$ shortest $s \sim u$ path

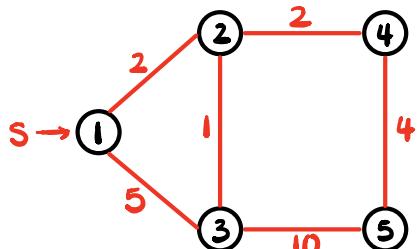
② min path $s \sim u$ via $X+u$:

$$\min(d[v], d[u] + l(u, v))$$

\downarrow
 $s \sim u - v$



	$X = \{1\}$	$X = \{1, 2\}$	$X = \{1, 2, 3\}$	$X = \{1, 2, 3, 4, 5\}$	$X = \{1, 2, 3, 4, 6\}$
$d[1]$	0	0	0	0	0
$d[2]$	2	2	2	2	2
$d[3]$	4	3	3	3	3
$d[4]$	∞	9	9	8	8
$d[5]$	∞	∞	6	6	6
$d[6]$	∞	∞	∞	11	9



	$X = \{1\}$	$X = \{1, 2\}$	$X = \{1, 2, 3\}$	$X = \{1, 2, 3, 4\}$
$d[1]$	0	0	0	0
$d[2]$	2	2	2	2
$d[3]$	5	3	3	3
$d[4]$	3	3	3	3
$d[5]$	∞	∞	13	7

Thm. unexplored u with min $d[u]$ value

$\Rightarrow d[u]$ is the $s \sim u$ shortest path

Pf. for any other path via ≥ 1 node z in $V - X$

$p: s \xrightarrow{\text{ex}} y - z \xrightarrow{\text{ex}} u$

$$l(p) = l(s \sim y) + l(y - z) + l(z \sim u)$$

$$\geq d[z]$$

$$d[u] \leq d[z] \Rightarrow d[u] \leq l(p)$$

$\Rightarrow d[u]$ is shortest path

Implementation push $d[u]$ & $u \in V$ into U

for $i = 1, \dots, n-1$

$u \leftarrow \min d[u] \text{ in } V - X \quad U.pop()$

for v s.t. $(u, v) \in E$

$d[v] = \min(d[v], d[u] + l(u, v)) \quad \text{update value in } U$

store d using array:

$O(n)$ to find min

$\Rightarrow O(n^2)$ total

store d using heap:

$O(n\log n + n\log n + n\log n)$

$\Rightarrow O(n\log n)$ total

another priority queue: fibonacci heap

$O(1)$ push

$O(\log n)$ pop

$O(1)$ value update

$$\Rightarrow O(n \log n + n + m) = O(n \log n + m) \text{ total}$$

4.5 Minimal Spanning Tree (MST)

• connected DG : $G = (V, E)$ with edge length

Goal: find a set of edges $T \subseteq E$

- $G' = (V, T)$ is connected

- minimize total cost $\sum_{e \in T} l(e)$

Sol. T^* will be a tree.

Find

Pf. (by γ)

If T^* has cycle C , assume $(u, v) \in C$.

$T^* - (u, v)$ is still a connected subgraph.

$\forall s, t$, if $s \sim u \sim v \sim t$ in T^*

$s \sim u \not\sim v \sim t$ in $T^* - (u, v)$

$T^* - (u, v)$ has smaller cost than T^* . γ

Prim's algorithm

maintain set X (current connected nodes)

$a[u]$ cost for adding u into X

$$X = \{s\}$$

$$T = \emptyset$$

$$a[u] = \begin{cases} l(s, u) & \text{if } (s, u) \in E \\ 0 & \text{otherwise} \end{cases}$$

for $i = 1, \dots, n-1$:

$u \leftarrow$ node in $V - X$ w/ smallest $a[u]$

add u to X

add $(\text{pre}[u], u)$ to T

for v s.t. $(u, v) \in E$

if $l(u, v) < a[v]$:

$$a[v] = l(u, v)$$

$$\text{pre}[v] = u$$

Time complexity: same as Dijkstra's algorithm

- heap for a : $O(m \log n)$

- fibonacci heap for a : $O(n \log n + m)$

$$T: (\text{pre}[u], u) \quad \forall u \in V$$

Pf. of correctness

Define "cut" node sets A, B :

$\text{cut}(A, B)$: all $(u, v) \in E$ s.t. $u \in A, v \in B$.

Prim's Alg.: $X, V-X$ add min-cost edge in $\text{cut}(A, B)$

Cut property: if edge e is the min-cost edge in $\text{cut}(X, V-X)$ to the node set X , then e must be in the MST.

Pf. If MST T^* and $e \notin T^*$:

assume $e = (u, v)$ and u, v in T^* are

connected by $\underbrace{u \sim p}_{\in X} - q \sim v$

Define $T' = T^* - (p, q) + (u, v)$

T' is still a spanning tree because $p \sim u - v \sim q$.

$l(T^*) > l(T')$. \downarrow

Kruskal's algorithm

- sort edges from small to large

- consider edges 1-by-1.

adding e doesn't make a cycle \Rightarrow add e
otherwise ignore e

Kruskal's algorithm

$$I = \emptyset$$

sort edges s.t. $e_1 \leq e_2 \leq \dots \leq e_m$

for $i = 1, \dots, m$:

$$(u, v) \leftarrow e_i$$

if u, v are in different connected components of T :

$$T = T + \{e_i\}$$

Pf. of correctness

If we add $e = (u, v)$ to T , u, v are in different connected components.

Assume X is the connected component of u .

e is the first edge in $\text{cut}(X, V - X)$ being considered in alg.

$\Rightarrow e$ is the min-cost edge in $\text{cut}(X, V - X)$

$\Rightarrow e$ should be in MST (**cut property**)

Implementation

- how to check u, v are in the same component?

Union-find structure: store n elements and their sets

- initial: each element forms individual set
- $\text{union}(u, v)$: merge sets of u, v
- $\text{find}(u)$: return set of u
 - $\text{find}(u) = \text{find}(v) \Rightarrow u, v$ in same set

Array implementation:

- find $O(1)$
- union $O(n)$

Array implementation (improved):

always rewrite smaller set

"Amortized complexity analysis": time complexity
for K unions from initial state

Claim: K unions take $O(K \log K)$ time

Pf. ① K unions \rightarrow touch at most $2K$ elements
② count how many times each element
is rewritten

rewrite to $u \Rightarrow$ size of set of u increases
by ≥ 2 times

assume p rewrites to u

\Rightarrow size of set of u is $\geq 2^p$

$2^p \leq 2K : p \leq \log(2K)$

\Rightarrow total rewrites = $O(K \log K)$