

Chapter 6 - Dynamic Programming

Lenny Wu

Professor Cho-Jui Hsieh

May 2021

Example: Fibonacci numbers.

$F_0 = 0, F_1 = 1, F_2 = 1, F_3 = 2, F_4 = 5, F_5 = 8, \dots$

Goal: compute n^{th} Fibonacci number ($n \geq 0$).

Recursive approach:

```
F(n):  
    if n == 0:  
        return 0  
    if n == 1:  
        return 1  
    return F(n - 1) + F(n - 2)
```

Problem with recursive approach: repeated function calls with the same input value; this wastes time and space (exponential time complexity). Instead, we can use dynamic programming with a top-down or bottom-up approach.

```
Fibonacci(n) (top-down):  
    M = [] // size n + 1  
    M[0] = 0, M[1] = 1:  
  
    if M[n] != empty:  
        return M[n]  
    else:  
        M[n] = F(n - 1) + F(n - 2)  
        return M[n]
```

```
Fibonacci(n) (bottom-up):  
    M = [] // size n + 1  
    M[0] = 0, M[1] = 1  
  
    for i = 2, 3, ..., n:  
        M[i] = M[i - 1] + M[i - 2]  
  
    return M[n]
```

Weighted Scheduling

Input: n jobs, each job has:

- s_n : starting time
- f_i : finishing time
- w_i : weight (value) of the job

Goal: find a set of compatible jobs S to maximize the total weight $\sum_{i \in S} w_i$.

DP: sort jobs by finishing time.

$$f_1 \leq f_2 \leq \dots \leq f_n$$

$OPT(i)$: optimal total weight solution for subproblem with jobs $\{1, \dots, i\}$.

The final solution is $OPT(n)$.

$$OPT(i) = \begin{cases} OPT(i-1) & \text{if job } i \text{ is not selected} \\ w_i + OPT(P_i) & \text{if job } i \text{ is selected} \end{cases}$$

P_i = max index such that P_i is not overlapping with i .

$$OPT(i) = \max(OPT(i-1), w_i + OPT(P_i)).$$

Algorithm:

DP (bottom up):

$M = []$ // size $n + 1$

$M[0] = 0$

for $i = 1, \dots, n$:

// $M[i] = OPT(i)$

$M[i] = \max(M[i-1], w_i + M[P_i])$

return $M[n]$

// computing P_i :

for each i :

find the last $j < i$ such that $f_j \leq s_i$ (\implies not overlapping)

\implies binary search ($f_1 \leq f_2 \leq \dots \leq f_{i-1}$)

Time complexity: binary search with $f_1 \leq f_2 \leq \dots \leq f_{i-1}$ is $O(\log n)$ for each i . Thus, the total is $O(n \log n)$.

Print the job list:

$A = []$ // size $n + 1$

$i = n$

while ($i > 0$):

```

if M[i - 1] < w_i + M[P_i]:
    A = A + {i}
    i = P_i
else:
    i -= 1

```

Time complexity: $O(n)$.

Knapsack: n items with values v_1, v_2, \dots, v_n , weights w_1, w_2, \dots, w_n .

find a subset S of items to maximize $\sum_{i \in S} v_i$ under the constraint $\sum_{i \in S} w_i \leq W_{max}$.

DP: $OPT(i, w)$: the solution with item $\{1, \dots, i\}$

$$OPT(i, w) = \begin{cases} v_i + OPT(i - 1, w - w_i) & \text{if choosing item } i \text{ (if } w \geq w_i) \\ OPT(i - 1, w) & \text{if not choosing } i \end{cases}$$

$$\longrightarrow OPT(i, w) = \max\{OPT(i - 1, w), v_i + OPT(i - 1, w - w_i)\}$$

Algorithm:

```

M = [] // size (n + 1) * (w_{max} + 1)
M[0][w] = 0 for all w

for i = 1, ..., n:
    for w = 0, ..., w_max:
        if w > w_i:
            M[i, w] = max(M[i - 1, w], v_i + M[i - 1, w - w_i])
        else:
            M[i, w] = M[i - 1, w]

return M[n, w_max]

```

Time complexity: $O(n \cdot W_{max})$ (not polynomial to input size; input value \neq size).

RNA Secondary Structure

Input: a sequence of $\{A, U, C, G\}$, and $B = b_1, \dots, b_n$.

Secondary structure is a set of pairs in this sequence which satisfies the following:

1. can only pair $(A, U), (U, A), (C, G), (G, C)$,
2. (non-sharp): pair (b_i, b_j) has to satisfy $i \leq j - 4$,
3. (non-crossing): for any two pairs $(b_i, b_j), (b_k, b_l)$, cannot have $i < k < j < l$.

Goal: find the maximum number of pairs that satisfy these 3 conditions.

1D DP (first attempt):

$OPT(i)$: optimal solution of b_1, \dots, b_i .

2D DP:

$OPT(i, j)$: solution (max number of pairs) in b_i, b_{i+1}, \dots, b_j .

Consider $\underbrace{b_i, b_{i+1}, \dots, b_{t-1}}_A \underbrace{b_t, \dots, b_j}_B$.

$$OPT(i, j) = \begin{cases} OPT(i, j-1) & \text{if } b_j \text{ is not in any pair} \\ 1 + OPT(i, t-1) + OPT(t+1, j-1) & \text{if } b_j \text{ is paired with some } b_t, i \leq t \leq j-4 \end{cases}$$

With b_t such that it can be paired with b_j , we have:

$$\longrightarrow OPT(i, j) = \max\{OPT(i, j-1), \max_{i \leq t \leq j-4} 1 + OPT(i, t-1) + OPT(t+1, j-1)\}$$

In paired case:

- 1 comes from (b_t, b_j) ,
- $OPT(i, t-1) = \#$ pairs in A ,
- $OPT(t+1, j-1) = \#$ pairs in B .

Algorithm:

```

OPT[i, j] = 0 when j <= i + 4    // initialize the table
for j = 1, ..., n:
  for i = 1, ..., j:
    compute OPT[i, j]    // for loop inside

```

Another ordering (increasing order of $|i - j|$):

```

for k = 5, ..., n - 1:
  for i = 1, ..., n - k:
    j = i + k
    compute OPT[i, j]

```

Time complexity: $O(n^3)$.

RNA Sequence Alignment

Input: two strings:

- $X = x_1, \dots, x_m$
- $Y = y_1, \dots, y_n$

Example:

$X = CUACCG$
 $Y = UACAUG$

Goal: insert, delete, and/or substitute to transform X into Y ; find the alignment with the smallest cost.

$$\text{Cost of each operation} = \begin{cases} \delta & \text{insert or delete} \\ \alpha_{x_i, y_j} & \text{substitute } x_i \text{ into } y_j \end{cases}$$

Worst case:

```

- - - - - C U A C C G
U A C A U G - - - - -

```

The alignment cost of this is $\delta \cdot (n + m)$.

Special case: no substitution; equivalent to longest common subsequence (LCS):

```

// in this example, LCS = 4:
C U A C C - - G
- U A C - A U G

```

The alignment cost of this is $\delta \cdot (n + m - 2 \cdot LCS)$.

2D DP:

$OPT(i, j) = \text{min alignment cost for } x_1, \dots, x_i \text{ and } y_1, \dots, y_j.$

The final solution is $OPT(m, n)$.

$$OPT(i, j) = \begin{cases} OPT(i-1, j-1) + \alpha_{x_i, y_j} & \text{match } (x_i, y_j) \\ OPT(i-1, j) + \delta & \text{delete } x_i \\ OPT(i, j-1) + \delta & \text{insert } y_i \text{ into } X \end{cases}$$

$$\longrightarrow OPT(i, j) = \min\{OPT(i-1, j-1) + \alpha_{x_i, y_j}, OPT(i-1, j) + \delta, OPT(i, j-1) + \delta\}$$

Algorithm:

```

M = [] // size (m + 1) * (n + 1)
M[i, 0] = i * delta for all i
M[0, i] = i * delta for all i

for i = 1, ..., m:

```

```

for j = 1, ..., n:
    M[i, j] = min( M[i - 1, j - 1] + alpha_{x_i, y_j},
                  M[i - 1, j] + delta,
                  M[i, j - 1] + delta )

```

```

return M[m, n]

```

Time complexity: $O(m \cdot n)$.

Space complexity: $O(m \cdot n)$. To reduce space to linear complexity: use two 1D arrays, one storing the current row and one storing the previous row.

How to get the optimal way of alignment?

Remember the previous node for each node.

- only one path from (m, n) to $(0, 0)$
- find shortest path from $(0, 0)$ to (m, n)

Run DP with linear space:

- $f(\frac{m}{2}, q)$ = shortest path from $(0, 0)$ to $(\frac{m}{2}, q)$,
- $g(\frac{m}{2}, q)$ = shortest path from $(\frac{m}{2}, q)$ to (m, n) .

How to find path with $O(m + n)$ space? Idea: divide and conquer:

- look at column $\frac{m}{2}$
- shortest path $(0, 0) \rightarrow (m, n)$ = shortest path $(0, 0) \rightarrow (\frac{m}{2}, q) \rightarrow (m, n)$

Algorithm:

```

find_path(X, Y):
    m = len(X), n = len(Y)

    // linear space DP; O(m*n) time
    compute f(m/2, q), g(m/2, q) for all q = 1, ..., n

    q* = argmin_q (f(m/2, q) + g(m/2, q))
    add (m/2, q*) to path

    find_path(X[1, ..., m/2], Y[1, ..., q*])    \ \ (1)
    find_path(X[m/2 + 1, ..., m], Y[q*, ..., n]) \ \ (2)

```

Recursive call (1) is $T(\frac{m}{2}, q)$, and recursive call (2) is $T(\frac{m}{2}, n - q)$.

→ $T(m, n) = T(\frac{m}{2}, q) + T(\frac{m}{2}, n - q) + c \cdot mn$.

→ $T(m, n) = O(mn)$.

Proof. $T(m, n) = O(mn)$.

Base case: $m = 1, n = 1$: $m + n = 1$ (trivial).

Inductive hypothesis: assume $T(m', n') \leq \alpha \cdot m'n'$ for all $m' + n' < m + n$.

Induction step:

$$\begin{aligned}T(m, n) &\leq T\left(\frac{m}{2}, q\right) + T\left(\frac{m}{2}, n - q\right) + c \cdot mn \\&\leq \alpha \cdot \frac{m}{2}q + \alpha \cdot \frac{m}{2}(n - q) + c \cdot mn \\&= \frac{\alpha}{2} \cdot mn + c \cdot mn \\&\leq \frac{\alpha}{2} \cdot mn + \frac{\alpha}{2} \cdot mn \\&= \alpha \cdot mn.\end{aligned}$$

Thus, $T(m, n) = O(mn)$.

□

Bellman-Ford

Goal: find the shortest path in a graph containing negative edges.

Why not use Dijkstra's algorithm?

1. May have negative cycles $\Rightarrow cost \rightarrow -\infty$.
2. Even without negative cycle, Dijkstra's algorithm does not work.

DP:

$OPT(i, v)$: min-cost path from v to t with $\leq i$ edges.

$$OPT(i, v) = \begin{cases} OPT(i-1, v) & \text{only use } i-1 \text{ edges: } v - w \rightsquigarrow t \\ OPT(i-1, w) + l(v, w) & \end{cases}$$

$$\rightarrow OPT(i, v) = \min\{OPT(i-1, v), \min_{w:(v,w) \in E} OPT(i-1, w) + l(v, w)\}$$

Algorithm:

```
M = [] // size n * |V|
for i = 1, ..., n - 1:
    for each node v:
        M[i, v] = min(M[i - 1, v], min(M[i - 1, w] + l(v, w)))
```

When should we stop?

Theorem. If a graph has no negative cycle, then there exists a shortest path with $\leq n - 1$ edges.

Proof. If shortest path has $\geq n$ edges, there exists a node visited twice in the path.

\Rightarrow There exists a cycle in the path.

No negative cycle \Rightarrow removing this cycle will not increase the cost. □

Algorithm using a single array:

```
M = [] // size n
for i = 1, ..., n - 1:
    for each node v:
        // M[v] is always <= M[i, v] at each iteration i
        M[v] = min(M[v], min(M[w] + l(v, w)))
```

Time complexity: $O(nm)$.

How to remember the shortest path? Remember the next node for each node:

```
next = [] // size |V|
for i = 1, ..., n - 1:
    for each node v:
        w* = argmin_w (M[w] + l(v, w))
```

$$\begin{aligned} \text{if } l(v, w^*) + M[w^*] < M[v] : \\ \quad \text{next}[v] &= w^* \\ M[v] &= l(v, w^*) \end{aligned}$$

What if there exists a negative cycle?

Theorem. *There exists a negative cycle that can reach t if and only if $OPT(n, v) < OPT(n-1, v)$ for some node v .*

Proof.

(\Rightarrow)

Suppose the graph has a negative cycle. Then there exists v such that $OPT(n, v) < OPT(n-1, v)$.

If $OPT(n, v) = OPT(n-1, v)$ for all v , then:

$$OPT(n+1, v) = OPT(n-1, v)$$

\vdots

$$OPT(i, v) = OPT(n-1, v) \text{ for all } i \geq n.$$

This contradicts with the definition of negative cycle since there exists v such that cost of path $v \rightsquigarrow t \rightarrow -\infty$.

(\Leftarrow)

Previous theorem: if there is no negative cycle, then there exists a shortest path with $\leq n-1$ edges. \square